Sachin Shastri

9 Garrison Avenue #44, Jersey City, NJ, 07306 | +1 (551) 358-7067 | sds662@nyu.edu
github.com/sachinsurfs | linkedin.com/in/sachindshastri

Education

New York University, Manhattan, New York, USA **Aug 2017 - May 2019**
Master of Science, Computer Science, Cum. GPA: 3.74 / 4.00
Courses: Distributed Systems, Predictive Analytics, Foundations of Machine Learning, Deep Learning, Natural Language Understanding, Computer Vision, Statistical NLP, Mathematics of Deep Learning, Fundamental Algorithms, Programming Languages, Operating Systems,

PES Institute of technology, Bangalore, India **Jun 2012 - Aug 2016**
Bachelor of Engineering, Computer Science
Cum. GPA: 9.04 / 10.0 - First class with Distinction (All Semesters)
Advanced courses: Big Data, Advanced C++, Applications of Machine Learning, Natural Language Processing

Skills

Languages: Over 10000 lines: Java • Python **Over 5000 lines**: C • C++ • Shell **Familiar**: Scala • R • Javascript
Machine learning aids: Keras • Tensorflow • PyTorch • Spark MLlib • SkLearn • Pandas
Automation/Build Tools: Jenkins • Docker • Kubernetes
Big-Data Frameworks: Hadoop • Storm • Spark • Beam • Kafka • Hive • Mahout • Oozie • Pig • Elasticsearch
Databases: MySQL • MongoDB • Cassandra • Teradata

Professional Experience

VIACOM | DATA SCIENCE INTERN **Jun 2018 – Aug 2018**
• Built an end-end machine learning pipeline that continuously learns from social media data and provides personalized strategies to brands to improve customer engagement and increase profit margin.
• Designed a ConvLSTM-based deep learning model to analyze user opinions on brands.

@WALMARTLABS | DATA ENGINEER **Jan 2016 – Aug 2017**
• Developed the first revisions of WM Datalake, Walmart's central data repository for the Analytics Division utilizing HDFS / S3.
• Built a data movement framework using Sqoop, Oozie & Hive for Mainframe / Teradata to fill the lake, used to migrate 1000+ daily jobs.
• Contributed in the complete development cycle of Walmart's in-house unified programming model for creating data processing pipelines with Spark, Storm, Hadoop, Kafka and Cassandra.
• Set up the cloud environment for continuous real-time processing 5k TPS(5GB/sec) of store transaction data.

Research Projects

Realtime autonomous traffic signal switching system | RL4J, Zigbee
A machine learning project, in collaboration with Bangalore Traffic Police, to optimize the traffic light switching system using reinforcement learning in a non-intrusive method to solve the traffic congestion problem in Bangalore.

Finding Connected components in Loss Surface | Tensorflow-GPU | Advisor: Prof. Joan Bruna
A study to investigate the nature of the loss surface of a deep neural network. Proposed an algorithmic technique to traverse different points of local minima from the hessian matrix decomposition of the loss function.

Let's Tell Tall Tales | Keras, PyTorch | Advisor: Prof. Sam Bowman
A study on different attention-based encoder-decoder RNN architectures to create robust language models that can produce coherent short stories when made to interact with each other while trained on different corpora.

Never Ending Learning of Sound | Weka, NumPy/SciPy, Tensorflow, OpenSmile, Audacity
A machine learning project, in association with Carnegie Mellon University, to design a web-based intelligent system that continuously searches and learns the meanings, associations, and semantics of sounds on the internet.

Academic Projects : • Contrastive Predictive Coding • Movie Memorable Quote Predictor • Cricket Match Simulation • LymphTracker • Multi-tasked Gender Predictor • Antibiotic Sensitivity Detector